

You have a few solid architectural options for integrating your vector search feature with your Tauri frontend + LLM tool backend. The best choice depends on factors like deployment, modularity, development workflow, and performance expectations.

Let me break down the trade-offs of each option:

---

## 1. Vector Search as an API (e.g., Actix-web)

### Pros:

- **Separation of concerns:** Your vector search service runs independently and exposes a clean HTTP/REST or gRPC API.
- **Language & Tech freedom:** You can build and deploy vector search in any stack without coupling it to Tauri.
- **Scalability & Reusability:** Other services or clients (not just Tauri) can call it easily.
- **Easier to scale/deploy separately:** You can scale vector search independently or deploy it on a dedicated server or cloud.
- **Better fault isolation:** Issues in vector search won't directly crash Tauri app.

### Cons:

- Adds complexity in deployment (need to run and manage another service).
  - Latency depends on IPC/HTTP calls (usually minimal locally, but must be considered).
- 

## 2. Include Vector Search Code Directly in the Tauri Backend

### Pros:

- Simple architecture, everything bundled together.
- Easier local integration with direct function calls — no IPC overhead.
- Single deployable binary for Tauri backend + vector search.

### Cons:

- Tightly couples vector search logic to your Tauri backend, harder to maintain or reuse.
  - Can lead to a monolithic codebase.
  - Vector search might have heavier dependencies that bloat your Tauri backend binary.
  - Less flexible scaling or separate upgrades.
-

### 3. Create a Workspace with Both Tauri and Vector Search as Separate Crates/Packages

#### Pros:

- Good middle ground — modular code but under one repo/workspace.
- You can share code (e.g., types, utils) between Tauri backend and vector search.
- Easier local development and testing — start/stop vector search as a separate process.
- Can still deploy vector search as a separate service (or embed it) depending on config.
- Cleaner project organization, better maintainability.

#### Cons:

- Slightly more complex than a single repo but manageable.
- Still need to decide on communication between Tauri backend and vector search (API calls, message passing, or direct calls if in the same binary).

---

## Recommendation

If your vector search is substantial, possibly resource-intensive, or might be reused by other services/clients, I'd recommend going with **option 1 (API/Actix-web)**. It's cleaner in the long term and better for scalability.

If you want simpler local integration and vector search is lightweight, **option 3 (workspace with separate crates)** gives you flexibility and maintainability, without fully committing to a separate service.

Avoid fully embedding vector search directly inside the Tauri backend (**option 2**) unless it's very small and tightly coupled, because it hurts modularity and maintainability.

---

## Extra thoughts:

- You could start with a workspace approach (option 3), build your vector search as a library or small binary, and expose a local HTTP server or an IPC mechanism.
- Then your Tauri backend calls it via HTTP or native Rust API.
- Later, you can separate it as a standalone API service if needed without major rewrites.

- For local apps, if vector search is heavy, running it as a background local service (API) is common.
-